

Assignment 1

Computer Architecture and Operating Systems

Name: Anmol Gupta
Roll Number: 2018329

Extract the kernel source twice to get two copies, and rename one of them. Here we will be making changes to `linux-3.16-rc-onlychanged`.

```
anmolgupta@ubuntu:~$  
anmolgupta@ubuntu:~$  
anmolgupta@ubuntu:~$  
anmolgupta@ubuntu:~$ ls  
linux-3.16-rc2  linux-3.16-rc2.tar.gz  linux-3.16-rc-onlychanged  
anmolgupta@ubuntu:~$ _
```

Create a folder `taskinfo`. This is where we will add the code for our system call.

```
linux-3.16-rc2  linux-3.16-rc2.tar.gz  linux-3.16-rc-onlychanged  
anmolgupta@ubuntu:~$ cd linux-3.16-rc-onlychanged/  
anmolgupta@ubuntu:~/linux-3.16-rc-onlychanged$ ls  
arch      crypto      fs          ipc         lib          net          scripts    tools  
block     Documentation  hello      Kbuild     MAINTAINERS  README      security   usr  
COPYING   drivers      include     Kconfig    Makefile     REPORTING-BUGS  sound      virt  
CREDITS   firmware     init        kernel     mm           samples      taskinfo  
anmolgupta@ubuntu:~/linux-3.16-rc-onlychanged$ cd taskinfo  
anmolgupta@ubuntu:~/linux-3.16-rc-onlychanged/taskinfo$ _
```

Create a file `taskinfo.h` and add the prototype of your system call here.

```
anmolgupta@ubuntu:~/linux-3.16-rc-onlychanged/taskinfo$ cat taskinfo.h  
asmlinkage long sys_sh_task_info(long pid, char *filename);  
  
anmolgupta@ubuntu:~/linux-3.16-rc-onlychanged/taskinfo$ _
```

Create a file `list_task_info.c` and write the system call in it.

```
#include<linux/kernel.h>
#include<linux/init.h>
#include<linux/sched.h>
#include<linux/syscalls.h>
#include<linux/file.h>
#include<asm/uaccess.h>
#include<uapi/asm-generic/errno-base.h>
#include<linux/fs.h>
#include<linux/fcntl.h>
```

We include the header files necessary in the system call. E.g., `sched.h` contains details about all the processes in the system. We also include `taskinfo.h` which contains our system call prototype.

```
#include "taskinfo.h"
```

```
int write_to_file(struct file *file, char *data) {
    file->f_op->write(file, data, strlen(data), &file->f_pos);
    return 0;
}
```

We use the write system call to write to the file. We supply the file descriptor, a character array and the length of data to be written.

```
asmlinkage long sys_sh_task_info(long pid, char *filename) {
```

```
    if(pid <= 0 || pid > 32768)
        return -EINVAL; // invalid argument
```

The maximum value of PID can be 32768. We check the PID and return if the argument is invalid.

```
    struct task_struct *task;
    struct file *file;
```

Create a pointer to the `task_struct` which we will use to get process details. Then create a file pointer which will represent an open file.

```
    char data[400], temp[400];
```

```
    loff_t pos = 0;
    int fd;
```

`loff_t` is used to determine the current position in the file.

```
    mm_segment_t old_fs = get_fs();
    set_fs(KERNEL_DS);
```

We increase the valid range of addresses for the buffer, by adding the kernel data space to it.

```
    fd = sys_open(filename, O_WRONLY|O_CREAT, 0644);
```

We open the file by specifying the filename and the mode (write only) in which we want to open it. We also create a file if it doesn't exist, and 0644 (owner can read and write) specifies the properties of the file if we create it.

System call code continued from the previous slide.

```
for_each_process(task) {  
    if(pid == (long) task->pid) {  
        printk("Process: %s\nPID_Number: %ld\nProcess State: %ld\nPriority: %ld\n",  
task->comm, (long) task->pid, (long) task->state, (long) task->prio);  
  
        sprintf(temp, "Process %s\nPID_Number: %ld\nProcess State: %ld\nPriority: %ld\n", task->comm, (long) task->pid, (long) task->state, (long) task->prio);  
  
        strcat(data, temp);  
  
        if(fd < 0)  
            return -EISDIR; // you cannot write to a directory  
  
        file = fget(fd);  
        write_to_file(file, data);  
    }  
}  
  
set_fs(old_fs);  
return 0;  
}
```

We use a macro defined in `sched.h` to iterate over all the processes.

We print some details about the required process to the kernel log. We also add it to the string `data` which we use to write to a file.

Check for unsuccessful opening of the file.

We change the limit of valid addresses back to the user address range.

We create a **Makefile** in the folder and add the following text to it. This is to make sure our system call is added to the kernel source code.

```
anmolgupta@ubuntu:~/linux-3.16-rc-onlychanged/taskinfo$ cat Makefile
obj-y:=list_task_info.o
anmolgupta@ubuntu:~/linux-3.16-rc-onlychanged/taskinfo$
```

We append the name of the directory **taskinfo** to this line in **Makefile** of the kernel source code. This tells the compiler where to look for our system call.

```
ifeq ($(KBUILD_EXTMOD),)
core-y      += kernel/ mm/ fs/ ipc/ security/ crypto/ block/ hello/ taskinfo/
```

Inside the source code folder, we go to **arch/x86/syscalls** and add the following line to the end of the **syscall_32.tbl** file.

```
352      i386      sched_getattr      sys_sched_getattr
353      i386      renameat2          sys_renameat2
354      i386      hello              sys_hello
355      i386      taskinfo           sys_sh_task_info
-- INSERT --
```

Inside the source code folder, we go to **include/linux** and add the new system call to the **syscalls.h** header file.

```
asmlinkage long sys_kcmp(pid_t pid1, pid_t pid2, int type,
                        unsigned long idx1, unsigned long idx2);
asmlinkage long sys_finit_module(int fd, const char __user *uargs, int flags);
asmlinkage long sys_sh_task_info(long pid, char *filename);_
#endif
```

We finally run the commands `make`, `make modules_install` and `make install` to compile the kernel source and integrate our system call.

We create a `test.c` file to test our system call and we'll store the output in the file `syscall_output`.

```
anmolgupta@ubuntu:~$ ls
linux-3.16-rc2  linux-3.16-rc2.tar.gz  linux-3.16-rc-onlychanged  syscall_output  test.c
anmolgupta@ubuntu:~$ _
```

Add code in `test.c` to test your system call.

```
#include<stdio.h>
#include<linux/kernel.h>
#include<sys/syscall.h>
#include<errno.h>
#include<unistd.h>

int main(int argc, char **argv) {
    int pid = atoi(argv[1]);
    long result = syscall(355, pid, argv[2]);
    if(result == 0)
        printf("System call sys_hello returned %ld\n", result);

    else {
        printf("Error :(\n");
        perror("Error ");
        printf("Error Number: %d\n", errno);
    }

    return 0;
}
```

We compile `test.c` and then run the executable `a.out` with two arguments--the **PID number of the process** and the **path of the file where process details have to be stored**. As expected, we see that the output gets written in the `syscall_output` file and is also printed to the kernel log.

```
anmolgupta@ubuntu:~$ gcc test.c
anmolgupta@ubuntu:~$ ./a.out 3 syscall_output
System call sys_hello returned 0
anmolgupta@ubuntu:~$ cat syscall_output
Process: ksoftirqd/0
PID Number: 3
Process State: 1
Priority: 120
anmolgupta@ubuntu:~$ _
```

We are getting the following details from the `task_struct`:

- i. **process name** - init, kthread, etc.
- ii. **process number** - unique process ID
- iii. **process state** - either created, ready, running, blocked or terminated state
- iv. **process priority** - used to carry out resource sharing between the processes

```
[ 29.723879] systemd-logind[1169]: New session c1 of user anmolgupta.
[10130.104247] Process: ksoftirqd/0
[10130.104250] PID Number: 3
[10130.104252] Process State: 1
[10130.104254] Priority: 120
anmolgupta@ubuntu:~$ history 2
1136 dmesg
1137 history 2
anmolgupta@ubuntu:~$ _
```

We have used `errno.h` header file to display the error messages for both **invalid argument due to a wrong PID** and trying to **write to a directory**.

```
anmolgupta@ubuntu:~$ ./a.out 0 syscall_output
Error :(
Error : Invalid argument
Error Number: 22
anmolgupta@ubuntu:~$ ./a.out 1 syscall_output/
Error :(
Error : Is a directory
Error Number: 21
anmolgupta@ubuntu:~$
```